

Problem Statement: This data set contains details of a Bank's customers and the target variable is a binary variable reflects on the fact whether the customer left the bank(closed his account) or not continues to be a customer.

In simple words, here I would **Predict whether the customer had left the organization or not.**

In the first portion, I have focussed on implementing a **basic neural network model** rather than concentrating on improving accuracy results.

```
Workflow

1.Import Necessary Libraries
2.Load the dataset
3.Drop the unnecessary columns
4.Check Null values and Missing values/NaN values
5.Use LabelEncoder to convert the Categorical data into Numerical data
6.Feature Selection and Data Splitting
7.Create Deep Learning Model:
    + 5 Steps to follow:
        1. Define the model architecture
        2. Compile the model
        3. Fit the model
        4. Evaluate the model
        5. Make Predictions

Here I have implemented the model MLP for Binary Classification

In a MultiLayer Perceptron model, the neural network is fully connected and comprised of layers of nodes where each node is connected to all the outputs from the previous layers and the output of each node is connected to all inputs for each of the nodes into the next layers.

importing tensorflow and necessary libraries

In [ ]:
import tensorflow as tf
from numpy import sqrt
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split

+ 5 Steps to follow:
    1. Define the model architecture
    2. Compile the model
    3. Fit the model
    4. Evaluate the model
    5. Make Predictions

In [ ]:
df.__version__ #checking the version

Out [ ]:
'2.5.0'

In [ ]:
from google.colab import files
uploaded = files.upload()

Loading Churn_Modelling.csv to Churn_Modelling.csv

Load the Dataset

In [ ]:
df = pd.read_csv('content/Churn_Modelling.csv')

Out [ ]:
   RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  EstimatedSalary  Exited
0            1    15634602  Hargrave      619         France  Female  42      2      0.00              1            1            1    101348.88      1
1            2    15647331         HBI       600          Spain  Female  41      1      0.000796             3            1            0    112542.58      0
2            3    15620004      Ouyi       502          France  Female  42      9    159660.80             3            1            0    113931.57      1
3            4    15707251      Buci       699          France  Female  39      1      0.00              2            0            0     93266.63      0
4            5    15737388      Michel      850          Spain  Female  43      2    125510.82             1            1            1     79044.10      0
...          ...          ...          ...          ...          ...          ...          ...          ...          ...          ...          ...          ...          ...
9995         9995    15606229  Oshijaku      771          France  Male  39      5      0.00              2            1            0     96270.64      0
9996         9997    15656892  Johansson     516          France  Male  35     10    57369.61             1            1            1    101699.77      0
9997         9999    15586532      Liu       709          France  Female  36      7      0.00              1            0            1     42095.58      1
9998         9999    15682955  Sabatini      772          Germany  Male  42      3    75075.31             2            1            0     92886.52      1
9999         9999    15620519  Walker      792          France  Female  28      4    130142.79             1            1            0     38190.78      0
10000 rows x 14 columns

In [ ]:
df.columns.tolist() #column list

Out [ ]:
['RowNumber',
 'CustomerId',
 'Surname',
 'CreditScore',
 'Geography',
 'Gender',
 'Age',
 'Tenure',
 'Balance',
 'NumOfProducts',
 'HasCrCard',
 'IsActiveMember',
 'EstimatedSalary',
 'Exited']

dropping unnecessary columns for prediction

In [ ]:
df = df.drop(['RowNumber', 'CustomerId', 'Surname'], axis=1)

Out [ ]:
   CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  EstimatedSalary  Exited
0            619         France  Female  42      2      0.00              1            1            1    101348.88      1
1            608          Spain  Female  41      1    83807.86             1            0            1    112542.58      0
2            502          France  Female  42      9    159660.80             3            1            0    113931.57      1
3            699          France  Female  39      1      0.00              2            0            0     93266.63      0
4            850          Spain  Female  43      2    125510.82             1            1            1     79044.10      0
...          ...          ...          ...          ...          ...          ...          ...          ...          ...          ...
9995         771          France  Male  39      5      0.00              2            1            0     96270.64      0
9996         516          France  Male  35     10    57369.61             1            1            1    101699.77      0
9997         709          France  Female  36      7      0.00              1            0            1     42095.58      1
9998         772          Germany  Male  42      3    75075.31             2            1            0     92886.52      1
9999         792          France  Female  28      4    130142.79             1            1            0     38190.78      0
10000 rows x 11 columns

checking null values

In [ ]:
df.isnull().sum()

Out [ ]:
CreditScore      0
Geography        0
Gender           0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64

checking NaN values

In [ ]:
df.isna().sum()

Out [ ]:
CreditScore      0
Geography        0
Gender           0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64

Checking the number of records and data type of the columns

In [ ]:
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
--  --
 0   CreditScore           10000 non-null  int64
 1   Geography             10000 non-null  object
 2   Gender               10000 non-null  object
 3   Age                 10000 non-null  int64
 4   Tenure              10000 non-null  int64
 5   Balance              10000 non-null  float64
 6   NumOfProducts        10000 non-null  int64
 7   HasCrCard            10000 non-null  int64
 8   IsActiveMember       10000 non-null  int64
 9   EstimatedSalary      10000 non-null  float64
10   Exited               10000 non-null  int64
dtypes: float64(2), int64(7), object(2)
memory usage: 659.5+ KB

Now, it is needed to convert the object type, categorical data to numerical data format.
```

```
Using Label Encoder to convert Categorical data to Numeric data

In [ ]:
#label_encoder = preprocessing.LabelEncoder()
#df['gender'] = label_encoder.fit_transform(df['gender'])
#df['geography'] = label_encoder.fit_transform(df['geography'])

In [ ]:
#label encoding for object type data
datatypes_dict = dict(df.dtypes)
#every check for mapping columns to labelencoder
LabelEncoderCollection = {}
for col_name, data_type in datatypes_dict.items():
    if data_type == 'object':
        LE = LabelEncoder()
        df[col_name] = LE.fit_transform(df[col_name])
        LabelEncoderCollection[col_name] = LE

In [ ]:
df.info() #all the data is in numeric form now

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
--  --
 0   CreditScore           10000 non-null  int64
 1   Geography             10000 non-null  int64
 2   Gender               10000 non-null  int64
 3   Age                 10000 non-null  int64
 4   Tenure              10000 non-null  int64
 5   Balance              10000 non-null  float64
 6   NumOfProducts        10000 non-null  int64
 7   HasCrCard            10000 non-null  int64
 8   IsActiveMember       10000 non-null  int64
 9   EstimatedSalary      10000 non-null  float64
10   Exited               10000 non-null  int64
dtypes: float64(2), int64(9)
memory usage: 859.5 KB
```

```
Now, it is needed to convert the object type, categorical data to numerical data format.

In [ ]:
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
--  --
 0   CreditScore           10000 non-null  int64
 1   Geography             10000 non-null  int64
 2   Gender               10000 non-null  int64
 3   Age                 10000 non-null  int64
 4   Tenure              10000 non-null  int64
 5   Balance              10000 non-null  float64
 6   NumOfProducts        10000 non-null  int64
 7   HasCrCard            10000 non-null  int64
 8   IsActiveMember       10000 non-null  int64
 9   EstimatedSalary      10000 non-null  float64
10   Exited               10000 non-null  int64
dtypes: float64(2), int64(9)
memory usage: 859.5 KB

In [ ]:
df

Out [ ]:
   CreditScore  Geography  Gender  Age  Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  EstimatedSalary  Exited
0            619         France  Female  42      2      0.00              1            1            1    101348.88      1
1            608          Spain  Female  41      1    83807.86             1            0            1    112542.58      0
2            502          France  Female  42      9    159660.80             3            1            0    113931.57      1
3            699          France  Female  39      1      0.00              2            0            0     93266.63      0
4            850          Spain  Female  43      2    125510.82             1            1            1     79044.10      0
...          ...          ...          ...          ...          ...          ...          ...          ...          ...          ...
9995         771          France  Male  39      5      0.00              2            1            0     96270.64      0
9996         516          France  Male  35     10    57369.61             1            1            1    101699.77      0
9997         709          France  Female  36      7      0.00              1            0            1     42095.58      1
9998         772          Germany  Male  42      3    75075.31             2            1            0     92886.52      1
9999         792          France  Female  28      4    130142.79             1            1            0     38190.78      0
10000 rows x 11 columns
```

```
Feature Selection and Data Splitting

In [ ]:
# split into input and output columns
X = df.values[:, :-1].astype('float32') #all the independent features
y = df.values[:, -1].astype(int) #splitted the target column, the dependent feature from the data

In [ ]:
# split into train and test datasets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(8000, 10) (2000, 10) (8000,) (2000,)

Creating Deep Learning Model

Step 1: Define the model

• In this step, the type of model building structure of the Deep Learning architecture can be decided and there are three steps- the Sequential Models, Functional API, or a custom architecture defined by the user.

• Depending on the problem, there are several architectures that can be used such as, CNN or ConvNets are used for computer vision tasks, then for natural language processing problems, RNN and LSTMs are preferable architectures.

• here we need to define the layers of the model, each layer will be configured with a number of nodes and activation function, and connecting the layers together.

In [ ]:
#start model with sequential object
model = tf.keras.models.Sequential()

#we are required to input object and specify the dimensions that you want to pass in
#dense layers means the hidden layers
model.add(tf.keras.layers.Dense(10)) #input layer shape=10, because the number of columns in final features are 10
#dense layers means the hidden layers
model.add(tf.keras.layers.Dense(32)) # the 2nd layer is 32 to increase complexities
# For the model evaluation procedure, some critical components of the training procedure is defined. Some necessary parameters need to be assigned in the following step such as, the loss, optimizers and the metrics.
# The model would be compiled through the optimizers, it is necessary to select the loss functions, optimization procedure(for example, stochastic gradient descent) or modern variations such as Adam, RMSProp, Adagrad or similar optimizers for computations can be used. The performance metrics are usually the accuracy or any user defined metrics for analysis to keep track during model training process.
# The optimizer can be specified as a string or instance can be created for an optimizer class, for example, 'sgd' for stochastic gradient descent.

• Usually these three loss functions are used -
    1.For binary classification : 'binary_crossentropy'
    2.For multi-class classification : 'sparse_categorical_crossentropy'
    3.For regression : 'mse' (mean squared error)

In [ ]:
#created optimizer and compile the model
#optimizer Adam has been used with learning rate

Optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer, loss='binary_crossentropy', metrics=['accuracy'])

Step 2: Fit the model

The third logical step is to fitting the model on the training dataset (model.fit()). The fit function trains the model for a fixed number of epochs. The term epochs means the iteration on a dataset.

• the important parameters such as the number of epochs, input and output data, validation data, the batch size (the number of samples in an epoch that estimates model error) for computing and calculating the essential features.

• a progress bar shows the training status of each epoch and the overall training process. Moreover, model performance can be simplified using the parameter 'verbose', when the value is 2. If verbose is set to 0, then summary will turn off.

In [ ]:
# fit the model / training the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=20)

# saves all the model metric performances while training, we can see the loss and accuracy of training and validation sets getting printed.

Epoch 1/20 ===== - 1s 2ms/step - loss: 1105.3628 - accuracy: 0.6526 - val_loss: 157.3726 - val_accuracy: 0.8920
Epoch 2/20 ===== - 1s 1ms/step - loss: 321.2197 - accuracy: 0.6670 - val_loss: 178.6638 - val_accuracy: 0.8920
Epoch 3/20 ===== - 0s 1ms/step - loss: 244.8924 - accuracy: 0.6729 - val_loss: 274.5328 - val_accuracy: 0.8920
Epoch 4/20 ===== - 0s 2ms/step - loss: 250.8243 - accuracy: 0.6814 - val_loss: 151.3423 - val_accuracy: 0.3620
Epoch 5/20 ===== - 0s 2ms/step - loss: 247.0703 - accuracy: 0.6861 - val_loss: 168.8297 - val_accuracy: 0.7115
Epoch 6/20 ===== - 0s 1ms/step - loss: 285.0220 - accuracy: 0.6731 - val_loss: 291.8254 - val_accuracy: 0.7655
Epoch 7/20 ===== - 0s 1ms/step - loss: 239.3802 - accuracy: 0.6754 - val_loss: 259.2396 - val_accuracy: 0.6535
Epoch 8/20 ===== - 0s 1ms/step - loss: 209.0990 - accuracy: 0.6895 - val_loss: 639.7608 - val_accuracy: 0.4390
Epoch 9/20 ===== - 0s 1ms/step - loss: 208.6722 - accuracy: 0.6730 - val_loss: 334.8481 - val_accuracy: 0.4810
Epoch 10/20 ===== - 0s 1ms/step - loss: 221.2949 - accuracy: 0.6731 - val_loss: 185.8244 - val_accuracy: 0.5215
Epoch 11/20 ===== - 0s 1ms/step - loss: 202.6965 - accuracy: 0.6780 - val_loss: 137.2355 - val_accuracy: 0.8020
Epoch 12/20 ===== - 0s 1ms/step - loss: 239.5986 - accuracy: 0.6760 - val_loss: 263.7321 - val_accuracy: 0.8020
Epoch 13/20 ===== - 0s 2ms/step - loss: 210.6725 - accuracy: 0.6861 - val_loss: 58.4352 - val_accuracy: 0.7888
Epoch 14/20 ===== - 0s 2ms/step - loss: 174.7944 - accuracy: 0.6889 - val_loss: 213.7567 - val_accuracy: 0.3955
Epoch 15/20 ===== - 0s 1ms/step - loss: 227.6772 - accuracy: 0.6755 - val_loss: 376.7634 - val_accuracy: 0.7820
Epoch 16/20 ===== - 0s 1ms/step - loss: 180.5550 - accuracy: 0.6759 - val_loss: 89.7110 - val_accuracy: 0.8518
Epoch 17/20 ===== - 0s 1ms/step - loss: 198.5553 - accuracy: 0.6796 - val_loss: 42.2615 - val_accuracy: 0.7785
Epoch 18/20 ===== - 0s 1ms/step - loss: 205.2193 - accuracy: 0.6791 - val_loss: 129.7933 - val_accuracy: 0.7530
Epoch 19/20 ===== - 0s 2ms/step - loss: 145.5674 - accuracy: 0.6754 - val_loss: 233.6974 - val_accuracy: 0.8020

#Accuracy Plot
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

#Loss Plot
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

model accuracy
model loss
```

Here, model learning is unstable as there are a lot of fluctuations in Accuracy Graph. The training accuracy is constant and the validation accuracy is highly fluctuating. A well-trained model would have training and validation accuracy going to gradually as epochs progress. The loss graph seems okay as they are decreasing gradually.

```
Step 4: Evaluate the model

• After training model evaluation (model.evaluate()) is done on the test dataset. The amount of data that is being used for prediction, has an impact on evaluation performance. The speed of model evaluation is proportional to the amount of your trained data. This is called the holdout category for model training.

Step 5: Make Predictions

• Apart from the trained dataset, the model's effectiveness is measured based on the prediction results on a random untrained datasets. Here simply calling the function (model.predict()) is used to predict the class label, probability or numerical values.
• Moreover, model evaluation metrics can be used such as - Classification Accuracy, Confusion matrix, Logarithmic Loss, Area under curve (AUC), F-Measure etc.

In [ ]:
# evaluate the model
from sklearn.metrics import accuracy_score, confusion_matrix
loss, acc = model.evaluate(x_test, y_test, verbose=0)
print('Test Accuracy: %.3f' % acc)

Test Accuracy: 0.802

In [ ]:
#collect predictions
predictions = np.round(model.predict(X_test))

In [ ]:
array([[0, 1],
       [0, 1],
       [0, 1],
       [0, 1],
       [0, 1], dtype=float32)

In [ ]:
#check the accuracy
from sklearn.metrics import accuracy_score, confusion_matrix
accuracy_score(y_test, predictions)

Out [ ]:
0.802

In [ ]:
confusion_matrix(y_test, predictions)

Out [ ]:
array([[1694,    0],
       [185,    6]])

Here model keeps predicting everything to be zero that is a common problem with imbalanced datasets where one class is very dominant over another class.

Now, I would be working on improving the accuracy using Dropout and Batch Normalization.
```

```
Improvement of Neural Network Model

Steps to improve Accuracy results in Neural Network:

• Add class weights to handle imbalance
• Increase the number of units and number of layers in Dense layers
• Add Batch Normalization to layers
• Add Dropout after every Input Layers

In [ ]:
#adding class weights to handle imbalance
from sklearn.utils.class_weight import compute_class_weight
class_weights = compute_class_weight(class_weight='balanced', classes=np.unique(y_train), y = y_train)

model.class_weights = {}
for e, weight in enumerate(class_weights):
    model.class_weights[e] = weight

model.class_weights

Out [ ]:
{0: 0.6290297216543481, 1: 2.437538886532602}

In [ ]:
#increase the number of units and number of layers in Dense layers
#add Batch Normalization to layers
#add Dropout after every Input Layers

model = tf.keras.models.Sequential()
model.add(tf.keras.layers.Dense(10))

model.add(tf.keras.layers.BatchNormalization()) #adding batch normalization before every layers
model.add(tf.keras.layers.Dense(128, activation='relu')) # value of neurons has been changed from 32 to 128, 64 respectively
model.add(tf.keras.layers.Dropout(0.02)) #deactivating 2% of neurons

model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dropout(0.02))

model.add(tf.keras.layers.BatchNormalization())
model.add(tf.keras.layers.Dense(32, activation='relu')) #before the final output layer, a dense layer with 32 neurons has been added.
model.add(tf.keras.layers.Dense(1, activation='sigmoid'))

model.summary()

WARNING:tensorflow:Please add 'keras.layers.InputLayer' instead of 'keras.Input' to Sequential model. 'keras.Input' is intended to be used by Functional mode
l.
Model: "sequential_1"

Layer (type)                Output Shape         Param #
=====
batch_normalization (BatchN (None, 10)           40
dense_3 (Dense)              (None, 128)          1488
dropout (Dropout)            (None, 128)           0
batch_normalization_1 (Batch (None, 128)          512
dense_4 (Dense)              (None, 64)           8256
dropout_1 (Dropout)          (None, 64)            0
batch_normalization_2 (Batch (None, 64)           256
dense_5 (Dense)              (None, 32)           2080
dense_6 (Dense)              (None, 1)             33
Total params: 12,585
Trainable params: 12,181
Non-trainable params: 404

Now the model has significantly higher parameters to train compared to the previous model.

In [ ]:
#compile the model
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer, loss='binary_crossentropy', metrics=['accuracy'])

In [ ]:
# fit the model / training the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=20, batch_size=32, verbose=2)

Epoch 1/20 - 1s - loss: 0.4390 - accuracy: 0.8105 - val_loss: 0.3881 - val_accuracy: 0.8455
Epoch 2/20 - 0s - loss: 0.3761 - accuracy: 0.8419 - val_loss: 0.3393 - val_accuracy: 0.8695
Epoch 3/20 - 0s - loss: 0.3598 - accuracy: 0.8481 - val_loss: 0.3298 - val_accuracy: 0.8695
Epoch 4/20 - 0s - loss: 0.3557 - accuracy: 0.8451 - val_loss: 0.3317 - val_accuracy: 0.8660
Epoch 5/20 - 0s - loss: 0.3522 - accuracy: 0.8515 - val_loss: 0.3250 - val_accuracy: 0.8670
Epoch 6/20 - 0s - loss: 0.3481 - accuracy: 0.8469 - val_loss: 0.3370 - val_accuracy: 0.8619
Epoch 7/20 - 0s - loss: 0.3472 - accuracy: 0.8544 - val_loss: 0.3284 - val_accuracy: 0.8695
Epoch 8/20 - 0s - loss: 0.3488 - accuracy: 0.8515 - val_loss: 0.3383 - val_accuracy: 0.8615
Epoch 9/20 - 0s - loss: 0.3469 - accuracy: 0.8522 - val_loss: 0.3298 - val_accuracy: 0.8660
Epoch 10/20 - 0s - loss: 0.3428 - accuracy: 0.8576 - val_loss: 0.3253 - val_accuracy: 0.8660
Epoch 11/20 - 0s - loss: 0.3423 - accuracy: 0.8598 - val_loss: 0.3289 - val_accuracy: 0.8645
Epoch 12/20 - 0s - loss: 0.3396 - accuracy: 0.8570 - val_loss: 0.3271 - val_accuracy: 0.8665
Epoch 13/20 - 0s - loss: 0.3374 - accuracy: 0.8587 - val_loss: 0.3264 - val_accuracy: 0.8655
Epoch 14/20 - 0s - loss: 0.3346 - accuracy: 0.8566 - val_loss: 0.3353 - val_accuracy: 0.8625
Epoch 15/20 - 0s - loss: 0.3349 - accuracy: 0.8594 - val_loss: 0.3311 - val_accuracy: 0.8625
Epoch 16/20 - 0s - loss: 0.3311 - accuracy: 0.8605 - val_loss: 0.3349 - val_accuracy: 0.8680
Epoch 17/20 - 0s - loss: 0.3277 - accuracy: 0.8662 - val_loss: 0.3423 - val_accuracy: 0.8660
Epoch 18/20 - 0s - loss: 0.3307 - accuracy: 0.8666 - val_loss: 0.3438 - val_accuracy: 0.8705
Epoch 19/20 - 0s - loss: 0.3282 - accuracy: 0.8583 - val_loss: 0.3432 - val_accuracy: 0.8678
Epoch 20/20 - 0s - loss: 0.3302 - accuracy: 0.8585 - val_loss: 0.3412 - val_accuracy: 0.8575

Now it takes less time for training compared to the previous model as there are more number of parameters.
```

```
In [ ]:
#accuracy Plot
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

#loss Plot
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(['train', 'validation'], loc='upper left')
plt.show()

model accuracy
model loss
```

Both the training and validation accuracies are gradually increasing now, and the loss kept decreasing with the number of epochs. This means that now the model has trained in quite well manner.

```
In [ ]:
#analyze the model
from sklearn.metrics import accuracy_score, confusion_matrix
loss, acc = model.evaluate(x_test, y_test, verbose=0)
print('Test Accuracy: %.3f' % acc)

Test Accuracy: 0.868

In [ ]:
#collect predictions
predictions = np.round(model.predict(X_test))

In [ ]:
#checking the accuracy
from sklearn.metrics import accuracy_score, confusion_matrix
accuracy_score(y_test, predictions)

Out [ ]:
0.8575

In [ ]:
confusion_matrix(y_test, predictions)

Out [ ]:
array([[1584, 189],
       [185, 211]])

The accuracy has increased a bit and through the confusion matrix, better prediction results are observed here than previous.

In [ ]:
from sklearn.metrics import classification_report
#the classification report visualizer displays the precision, recall, F1, and support scores for the model.
print(classification_report(y_test, predictions))

precision    recall  f1-score   support

accuracy      0.86      0.86      0.86      2000
macro avg      0.79      0.74      0.76      2000
weighted avg   0.85      0.86      0.85      2000
```